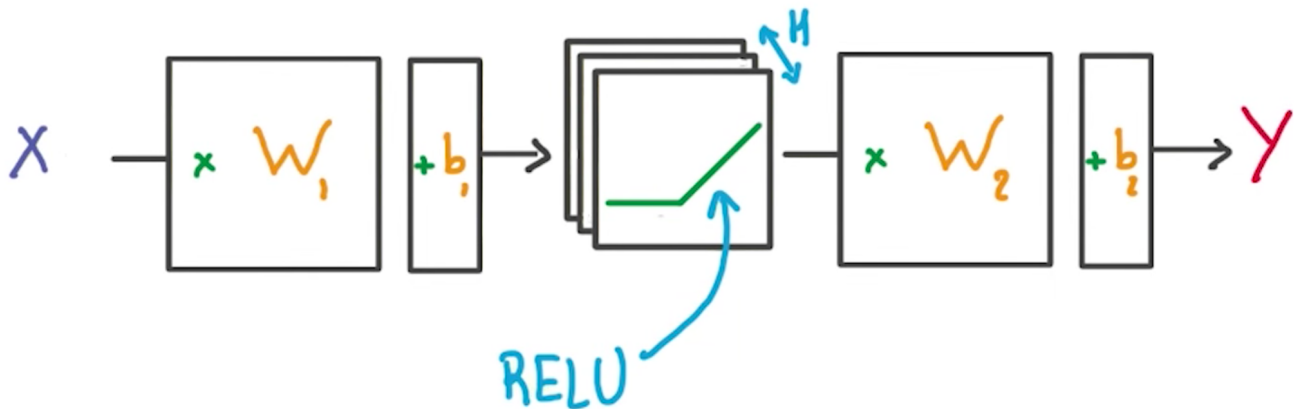


`tf.reshape()` function above reshapes the 28px by 28px matrices in `x` into row vectors of 784px.

Multilayer Perceptron



```
# Hidden layer with RELU activation
layer_1 = tf.add(tf.matmul(x_flat, weights['hidden_layer']),\
    biases['hidden_layer'])
layer_1 = tf.nn.relu(layer_1)
# Output layer with linear activation
logits = tf.add(tf.matmul(layer_1, weights['out']), biases['out'])
```

You've seen the linear function

```
tf.add(tf.matmul(x_flat, weights['hidden_layer']), biases['hidden_layer'])
```

before, also known as $xw + b$. Combining linear functions together using a ReLU will give you a two layer network.

Optimizer

```
# Define loss and optimizer
cost = tf.reduce_mean(\
    tf.nn.softmax_cross_entropy_with_logits(logits=logits, labels=y))
optimizer = tf.train.GradientDescentOptimizer(learning_rate=learning_rate)\
    .minimize(cost)
```

This is the same optimization technique used in the Intro to TensorFlow lab.