


Exploration vs. Exploitation



Exploration-Exploitation Dilemma ([Source](#))

Solving Environments in OpenAI Gym

In many cases, we would like our reinforcement learning (RL) agents to learn to maximize reward as quickly as possible. This can be seen in many OpenAI Gym environments.

For instance, the [FrozenLake-v0](#) environment is considered solved once the agent attains an average reward of 0.78 over 100 consecutive trials.